

```

import os
import base64
import numpy as np
import tensorflow as tf

# Get current working directory
current_dir = os.getcwd()

# Append data/mnist.npz to the previous path to get the full path
data_path = os.path.join(current_dir, "/content/DL2.html")

# Load data (discard test set)
(training_images, training_labels), _ = tf.keras.datasets.mnist.load_data(path=dat

print(f"training_images is of type {type(training_images)}.\ntraining_labels is of

# Inspect shape of the data
data_shape = training_images.shape

print(f"There are {data_shape[0]} examples with shape ({data_shape[1]}, {data_shap

⇒ training_images is of type <class 'numpy.ndarray'>.
   training_labels is of type <class 'numpy.ndarray'>

   There are 60000 examples with shape (28, 28)

def reshape_and_normalize(images):
    """Reshapes the array of images and normalizes pixel values.

    Args:
        images (numpy.ndarray): The images encoded as numpy arrays

    Returns:
        numpy.ndarray: The reshaped and normalized images.
    """

    ### START CODE HERE ###

    # Reshape the images to add an extra dimension (at the right-most side of the
    images = np.expand_dims(images, axis=-1)

    # Normalize pixel values
    images = images/255.0

    ### END CODE HERE ###

    return images

# Reload the images in case you run this cell multiple times
(training_images, _), _ = tf.keras.datasets.mnist.load_data(path=data_path)

```

```
# Apply your function
training_images = reshape_and_normalize(training_images)

print('Name: A.ARUVI.          RegisterNumber: 21222230014.          \n')
print(f"Maximum pixel value after normalization: {np.max(training_images)}\n")
print(f"Shape of training set after reshaping: {training_images.shape}\n")
print(f"Shape of one image after reshaping: {training_images[0].shape}")
```

```
➡ Name: A.ARUVI.          RegisterNumber: 21222230014.

Maximum pixel value after normalization: 1.0

Shape of training set after reshaping: (60000, 28, 28, 1)

Shape of one image after reshaping: (28, 28, 1)
```

```
# GRADED CLASS: EarlyStoppingCallback
```

```
### START CODE HERE ###
```

```
# Remember to inherit from the correct class
```

```
class EarlyStoppingCallback(tf.keras.callbacks.Callback):
```

```
    # Define the correct function signature for on_epoch_end method
```

```
    def on_epoch_end(self, epoch, logs={}):
```

```
        # Check if the accuracy is greater or equal to 0.98
```

```
        if (logs.get('accuracy')>=0.995):
```

```
            # Stop training once the above condition is met
```

```
            self.model.stop_training=True
```

```
            print("\nReached 98% accuracy so cancelling training!")
```

```
print('Name: A.ARUVI.          Register Number: 21222230014.          \n')
```

```
### END CODE HERE ###
```

```
➡ Name: A.ARUVI.          Register Number: 21222230014.
```

```
def convolutional_model():
```

```
    """Returns the compiled (but untrained) convolutional model.
```

```
    Returns:
```

```
        tf.keras.Model: The model which should implement convolutions.
```

```
    """
```

```
## START CODE HERE ##
```

```
# Define the model
```

```
model = tf.keras.models.Sequential([
```

```
    # Add convolutions and max pooling
```

```
    tf.keras.Input(shape=(28,28,1)),
```

```
    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
```

```

tf.keras.layers.MaxPooling2D(2, 2),
tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
tf.keras.layers.MaxPooling2D(2,2),

# Add the same layers as before
tf.keras.layers.Flatten(),
tf.keras.layers.Dense(128, activation='relu'),
tf.keras.layers.Dense(10, activation='softmax')
])

```

```

### END CODE HERE ###

```

```

# Compile the model
model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)
return model

```

```

model.summary()

```

➞ Model: "sequential_1"

Layer (type)	Output Shape	
conv2d_2 (Conv2D)	(None, 26, 26, 64)	
max_pooling2d_2 (MaxPooling2D)	(None, 13, 13, 64)	
conv2d_3 (Conv2D)	(None, 11, 11, 64)	
max_pooling2d_3 (MaxPooling2D)	(None, 5, 5, 64)	
flatten_1 (Flatten)	(None, 1600)	
dense_2 (Dense)	(None, 128)	
dense_3 (Dense)	(None, 10)	

Total params: 731,360 (2.79 MB)
 Trainable params: 243,786 (952.29 KB)
 Non-trainable params: 0 (0.00 B)
 Optimizer params: 487,574 (1.86 MB)

```

### Model compiling and Training

```

```

# Define your compiled (but untrained) model
model = convolutional_model()


```

```

# Train your model (this can take up to 5 minutes)
training_history = model.fit(training_images, training_labels, epochs=10, callbacks=[Earl

```

➞ Epoch 1/10
 1875/1875 ————— 99s 52ms/step - accuracy: 0.9109 - loss: 0.2842
 Epoch 2/10

1875/1875  **90s** 48ms/step - accuracy: 0.9860 - loss: 0.0433
Epoch 3/10
1875/1875  **91s** 48ms/step - accuracy: 0.9919 - loss: 0.0262
Epoch 4/10
1875/1875  **146s** 51ms/step - accuracy: 0.9943 - loss: 0.0178
Epoch 5/10
1875/1875  **0s** 48ms/step - accuracy: 0.9957 - loss: 0.0134
Reached 98% accuracy so cancelling training!
1875/1875  **91s** 48ms/step - accuracy: 0.9957 - loss: 0.0134

Start coding or [generate](#) with AI.